



# Big Data Overview



## Section Overview

- Explanation of Hadoop, MapReduce, Spark, and PySpark
- Local versus Distributed Systems
- Overview of Hadoop Ecosystem
- Detailed overview of Spark
- Set-up on Amazon Web Services
- Resources on other Spark Options
- Jupyter Notebook hands-on code with PySpark and RDDs



# Big Data


- We've worked with Data that can fit on a local computer, in the scale of 0-8 GB.
- But what can we do if we have a larger set of data?
  - Try using a SQL database to move storage onto hard drive instead of RAM
  - Or use a distributed system, that distributes the data to multiple machines/computer.



# Local versus Distributed




Local




Distributed



# Local versus Distributed



Local



Distributed



# Big Data

- A local process will use the computation resources of a single machine
- A distributed process has access to the computational resources across a number of machines connected through a network
- After a certain point, it is easier to scale out to many lower CPU machines, than to try to scale up to a single machine with high a CPU



# Big Data

- Distributed machines also have the advantage of easily scaling, you can just add more machines
- They also include fault tolerance, if one machine fails, the whole network can still go on.
- Let's discuss the typical format of a distributed architecture that uses Hadoop




- Hadoop is a way to distribute very large files across multiple machines.
- It uses the Hadoop Distributed File System (HDFS)
- HDFS allows a user to work with large data sets
- HDFS also duplicates blocks of data for fault tolerance
- It also then uses MapReduce
- MapReduce allows computations on that data






# Distributed Storage - HDFS





# Distributed Storage - HDFS


- HDFS will use blocks of data, with a size of 128 MB by default
- Each of these blocks is replicated 3 times
- The blocks are distributed in a way to support fault tolerance





# Distributed Storage - HDFS

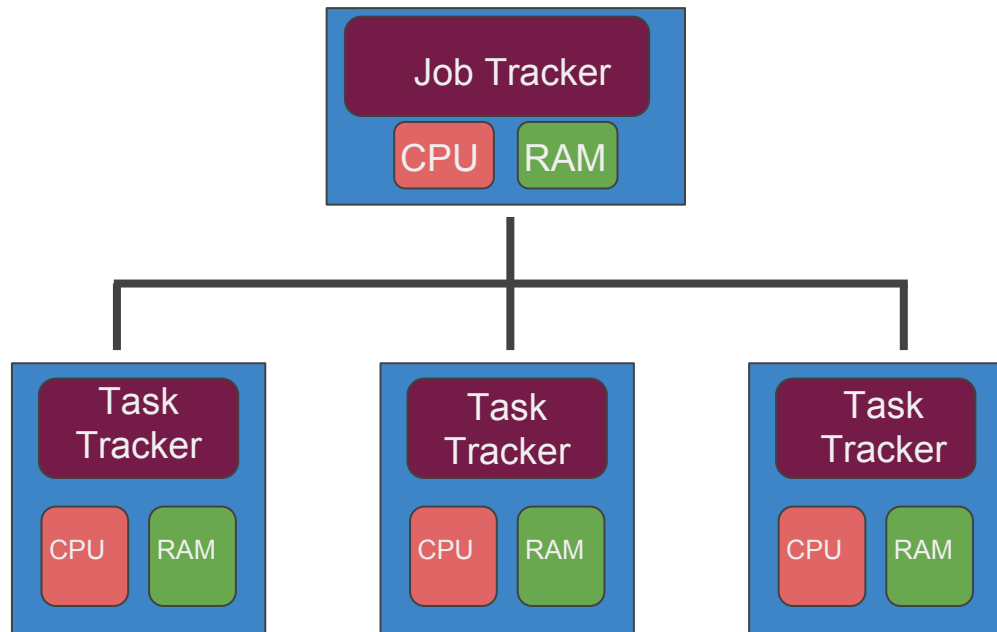
- Smaller blocks provide more parallelization during processing
- Multiple copies of a block prevent loss of data due to a failure of a node





# MapReduce


- MapReduce is a way of splitting a computation task to a distributed set of files (such as HDFS)
- It consists of a Job Tracker and multiple Task Trackers





# MapReduce

- The Job Tracker sends code to run on the Task Trackers
- The Task trackers allocate CPU and memory for the tasks and monitor the tasks on the worker nodes





# Big Data

- What we covered can be thought of in two distinct parts:
  - Using HDFS to distribute large data sets
  - Using MapReduce to distribute a computational task to a distributed data set
- Next we will learn about the latest technology in this space known as Spark.
- Spark improves on the concepts of using distribution



# Up next we'll learn about Spark!



# Spark Overview





- This lecture will be an abstract overview, we will discuss:
  - Spark
  - Spark vs MapReduce
  - Spark RDDs
  - RDD Operations
- Don't worry about having to understand all the operations, we will review and cover this again when we actually program with Spark and Python



- Spark is one of the latest technologies being used to quickly and easily handle Big Data
- It is an open source project on Apache
- It was first released in February 2013 and has exploded in popularity due to its ease of use and speed
- It was created at the AMPLab at UC Berkeley



- You can think of Spark as a flexible alternative to MapReduce
- Spark can use data stored in a variety of formats
  - Cassandra
  - AWS S3
  - HDFS
  - And more



# Spark vs MapReduce

- MapReduce requires files to be stored in HDFS, Spark does not!
- Spark also can perform operations up to 100x faster than MapReduce
- So how does it achieve this speed?



# Spark vs MapReduce

- MapReduce writes most data to disk after each map and reduce operation
- Spark keeps most of the data in memory after each transformation
- Spark can spill over to disk if the memory is filled




# Spark RDDs

- At the core of Spark is the idea of a Resilient Distributed Dataset (RDD)
- Resilient Distributed Dataset (RDD) has 4 main features:
  - Distributed Collection of Data
  - Fault-tolerant
  - Parallel operation - partitioned
  - Ability to use many data sources

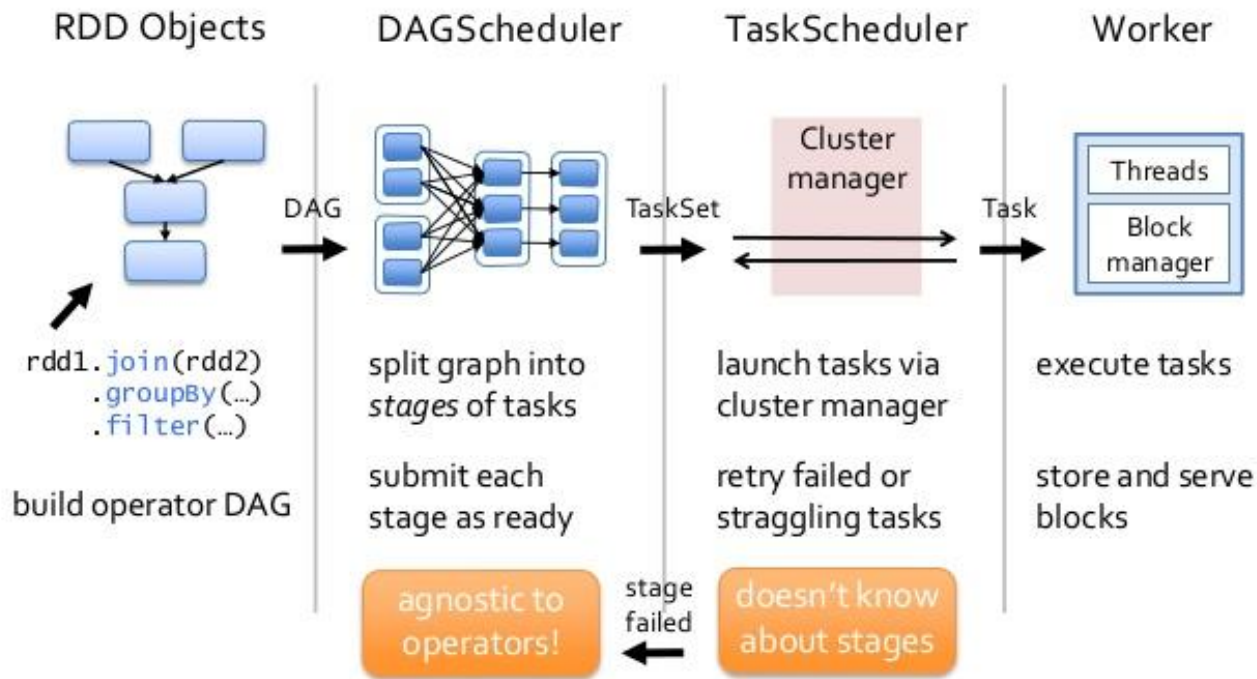


# Spark RDDs





# Spark RDDs







# Spark RDDs

- RDDs are immutable, lazily evaluated, and cacheable
- There are two types of RDD operations:
  - Transformations
  - Actions



# Spark RDDs

- Basic Actions
  - First
  - Collect
  - Count
  - Take



# Spark RDDs

- Collect - Return all the elements of the RDD as an array at the driver program.
- Count - Return the number of elements in the RDD
- First - Return the first element in the RDD
- Take - Return an array with the first n elements of the RDD




# Spark RDDs

- Basic Transformations
  - Filter
  - Map
  - FlatMap



# Spark RDDs


- `RDD.filter()`
  - Applies a function to each element and returns elements that evaluate to true





# Spark RDDs


- `RDD.map()`
  - Transforms each element and preserves # of elements, very similar idea to pandas `.apply()`





# Spark RDDs

- `RDD.flatMap()`
  - Transforms each element into 0-N elements and changes # of elements





# Map vs FlatMap

- Map()
  - Grabbing first letter of a list of names
- FlatMap()
  - Transforming a corpus of text into a list of words
- We will show many more examples when programming with PySpark





## Pair RDDs

- Often RDDs will be holding their values in tuples
  - (key,value)
- This offers better partitioning of data and leads to functionality based on reduction



# Reduce and ReduceByKey

- `Reduce()`
  - An action that will aggregate RDD elements using a function that returns a single element
- `ReduceByKey()`
  - An action that will aggregate Pair RDD elements using a function that returns a Pair RDD
- These ideas are similar to a Group By operation



# Big Data

- Spark is being continually developed and new releases come out often!
- The Spark Ecosystem now includes:
  - Spark SQL
  - Spark DataFrames
  - MLlib
  - GraphX
  - Spark Streaming



# Big Data

- Now we've learned enough to get started!
- We're now going to show you how to set up an Amazon Web Services account to get Spark up and running!
- We'll also have text article lecture for some other options in case you don't want to use AWS.



# AWS Account Set-Up



- Go to:
  - <https://aws.amazon.com/free>
- Then click on Create Free Account
- Sign up with an email address
- Then fill out the profile information



- The profile information:
  - Contact Information
  - Billing Information
  - ID Verification
  - Choose free support plan
- Next lecture we will explore AWS and create an EC2 instance



# EC2 Instance Set-up






- Now that we have our AWS account we will create an EC2 Instance.
- Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides resizable compute capacity in the cloud.
- We can basically think of it as a virtual computer we can access through the internet.




- So here is our plan:
  - Create EC2 Instance on AWS
  - Use SSH to connect to EC2 over internet
  - SSH is different for Windows vs Mac/Linux
  - Set-up Spark and Jupyter on EC2 Instance





- SSH (Secure Shell Connection)
  - Watch this lecture all the way through for Windows
  - Skip to next lecture after EC2 Set-up for Mac/Linux
- Our goal is to remotely connect to the command line of our virtual machine on Amazon EC2





**Login to your AWS console at:  
[aws.amazon.com](https://aws.amazon.com)**



# PySpark Set-up



# SSH with Mac and Linux



- Skip this lecture if you are on Windows, you should have connected to your instance already from the previous lecture
- We've created our EC2 instance using AWS console
- You should have also downloaded the .pem file
- Now we are going to connect to our instance through our terminal using SSH



- Make sure you can locate your .pem file
  - Recommend you relocate it to your Desktop
- Make sure you have the DNS address of your EC2 instance
- Check the resource link for the step by step instructions from Amazon
- Let's get started by opening our terminal



